Unit 6: Simple linear regression
1. Introduction to regression

Sta 104 - Summer 2015

Duke University, Department of Statistical Science

June 12, 2015

Dr. Çetinkaya-Rundel

Slides posted at http://bit.ly/sta104su15

Clicker question

Which of the following is the best guess for the correlation between annual murders per million and percentage living in poverty?

(a) -1.52
(b) -0.63
(c) -0.12
(d) 0.02
(e) 0.84

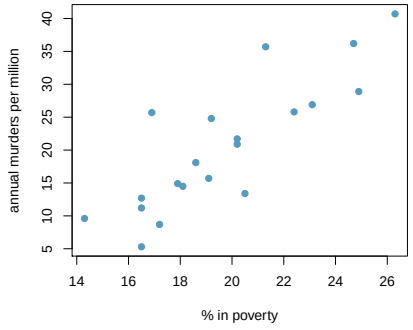

Clicker question

Which of the following is the best guess for the correlation between annual murders per million and population size?

(a) -0.97
(b) -0.61
(c) -0.06
(d) 0.55
(e) 0.97

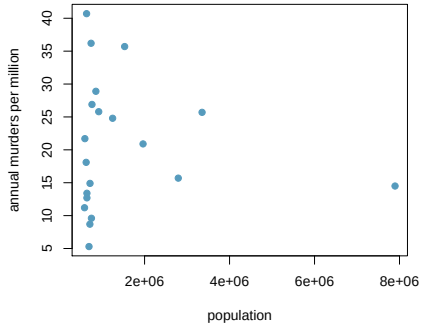

Assessing the correlation

Clicker question

Which of the following is has the strongest correlation, i.e. correlation coefficient closest to +1 or -1?

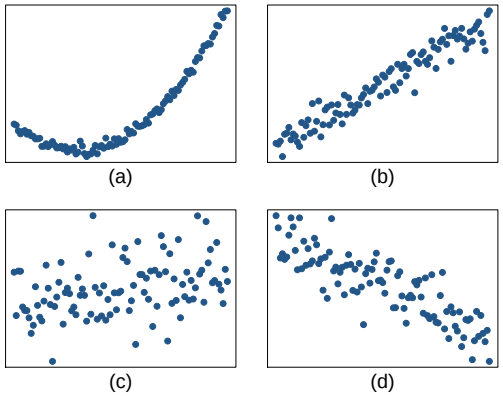

Play the game!

Post a screenshot with your problem set HW for +1 pt extra credit!
http://mih5.github.io/correlation_game2/correlationgame.html

Spurious correlations

Remember: correlation does not always imply causation!
http://www.tylervigen.com/

(2) Least squares line minimizes squared residuals

- Residuals are the leftovers from the model fit, and calculated as the difference between the observed and predicted y: $e_i = y_i - \hat{y}_i$
- The least squares line minimizes squared residuals:
  - Population data: $\hat{y} = \beta_0 + \beta_1 x$
  - Sample data: $\hat{y} = b_0 + b_1 x$

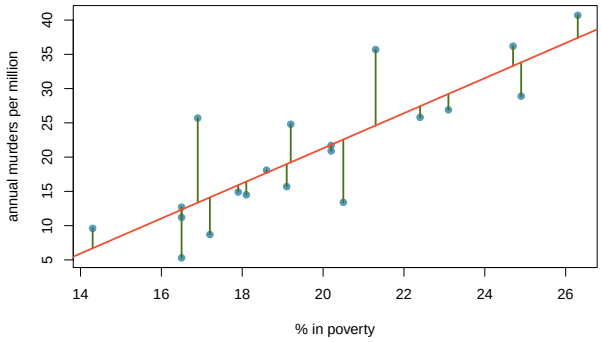

(3) Interpreting the last squares line

- *Slope:* For each unit increase in x, y is expected to behigher/lower on average by the slope.

$$b_1 = \frac{s_y}{s_x} R$$

- *Intercept:* When $x = 0$, y is expected to equal the intercept.

$$b_0 = \bar{y} - b_1 \bar{x}$$

  - The calculation of the intercept uses the fact the a regression line **always** passes through $(\bar{x}, \bar{y})$.

Why does the regression line **always** pass through $(\bar{x}, \bar{y})$?

- If there is no relationship between x and y ($b_1 = 0$), the best guess for $\hat{y}$ for any value of x is $\bar{y}$.
- Even when there is a relationship between x and y ($b_1 \neq 0$), the best guess for $\hat{y}$ when $x = \bar{x}$ is still $\bar{y}$.

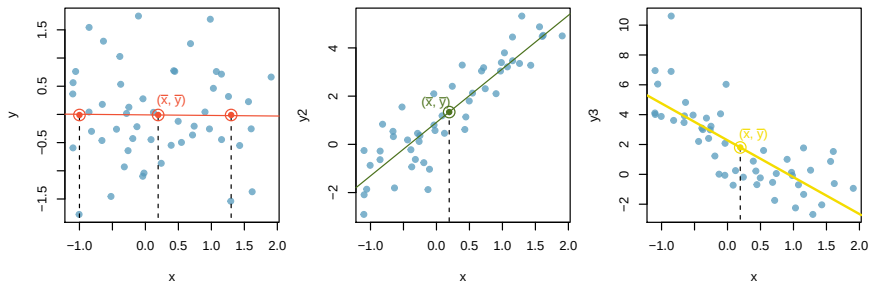

Application exercise: 6.1 Linear model

See course website for details

Clicker question

What is the interpretation of the slope?

(a) Each additional percentage in those living in poverty increases number of annual murders per million by 2.56.

(b) For each percentage increase in those living in poverty, the number of annual murders per million is expected to be higher by 2.56 on average.

(c) For each percentage increase in those living in poverty, the number of annual murders per million is expected to be lower by 29.91 on average.

(d) For each percentage increase annual murders per million, the percentage of those living in poverty is expected to be higher by 2.56 on average.

Clicker question

Suppose you want to predict annual murder count (per million) for a series of districts that were not included in the dataset. For which of the following districts would you be most comfortable with your prediction?

A district where % in poverty =

(a) 5%
(b) 15%
(c) 20%
(d) 26%
(e) 40%

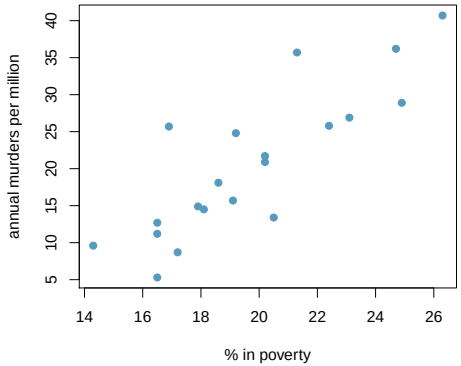

A note about the intercept

Sometimes the intercept might be an extrapolation: useful for adjusting the height of the line, but meaningless in the context of the data.

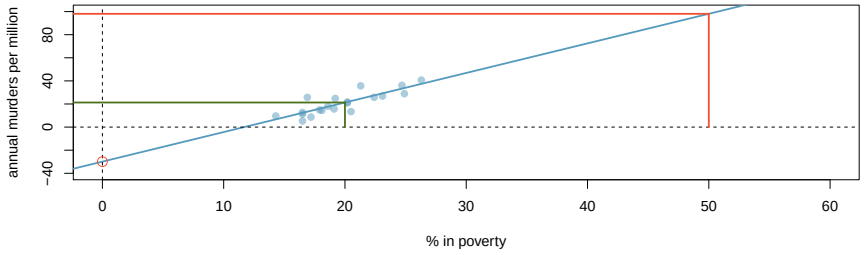

Summary of main ideas

1. Correlation coefficient describes the strength and direction of the linear association between two numerical variables
2. Least squares line minimizes squared residuals
3. Interpreting the last squares line
4. Predict, but don't extrapolate